MunBench: A Generative Benchmark for Korean Emotional Intelligence, Creative Writing, and Cultural Nuance in Large Language Models

Taewook Ha
hataewook12@gmail.com
https://github.com/ghandhitechnology/munbench

July 2026

Abstract

Korean-language evaluation of large language models (LLMs) is dominated by multiple-choice knowledge tests. These are cheap to grade and resistant to ambiguity, but they cannot answer the question practitioners actually ask: *can this model write well in Korean?* Prior work shows that factual knowledge of a culture correlates only weakly with the ability to apply that culture in open-ended generation, so multiple-choice scores are a poor proxy. We present **MunBench** (문벤치), a fully generative benchmark that evaluates three abilities at once: (1) *emotional intelligence* in Korean interpersonal context, tested through multi-turn roleplay scenarios built around 눈치 (reading the unstated), 정 (the deep affective bond formed over shared time, with its attendant obligations), and 체면 (face); (2) *creative writing*, tested through 48 commissions that each target a documented LLM weakness such as melodrama and cliché; and (3) *cultural and linguistic nuance*, tested through honorific-register switching, idiom deployment, sarcasm resolution, and paired prompts that reveal whether a model silently defaults to Western norms when a Korean setting is implied but not stated. All 158 items were authored natively in Korean. Scoring combines an ensemble rubric, anchored pairwise Elo, and judge-free mechanical metrics that catch failure modes LLM judges are known to miss, including code-switching into English. In an initial 13-model study, three results stand out. The creative-writing track separates models that other tracks cannot: writing scores span 1.7–8.2, while eight of thirteen models sit within one point of each other on the EQ track. The culture-pair probe finds four models whose scores *drop* when a Korean setting is made explicit — by 2.8 rubric points in the worst case — a pattern no culture-specified benchmark can see. And a judge-free tripwire catches a frontier model code-switching out of Korean in 42% of nuance-track outputs, a failure its rubric scores alone would blur. Benchmark, harness, and results are public.

1 Introduction

A user who asks an LLM for help in Korean is rarely asking it to answer exam questions. They are asking it to write: a reply to a delicate message, a eulogy, a cover letter, a story. Whether the model does this *well* in Korean — with the right honorific register, without translated-English phrasing, reading what the other person meant rather than what they said — is the ability that matters in practice, and it is almost entirely unmeasured.

Korean LLM benchmarks overwhelmingly use multiple-choice formats: KMMLU [Son et al., 2024b], HAE-RAE Bench [Son et al., 2023], CLIcK [Kim et al., 2024a], KoBEST [Jang et al., 2022], and KorNAT [Lee et al., 2024] all grade selection, not production. This is a reasonable design choice for knowledge, but it leaves generation quality unmeasured, and there is direct evidence that the two dissociate: Nunchi-Bench [Kim and Lee, 2025] found that accuracy on factual questions about Korean cultural practices correlates only weakly with appropriate behavior in open-ended scenarios that depend on those same practices. Knowing *about* 눈치 is not the same as having it. The few Korean benchmarks that do score generation have used a single generic quality scale; LogicKor, the most widely used, was retired after frontier models saturated its 0–10 judge rubric. Meanwhile, the English-language evaluation community has developed strong methodology for exactly this problem — EQ-Bench 3's judged roleplay with paired rubric-and-Elo scoring [Paech, 2023], bias controls for LLM judges [Zheng et al., 2023], and reward-model findings from LitBench [Fein et al., 2025] — but none of it has been applied to Korean.

MunBench fills this gap. It is, to our knowledge, the first benchmark that evaluates Korean *generation* across emotional intelligence, creative writing, and cultural nuance simultaneously, with methodology adapted from the strongest English-language precedents and grounded in Korean-specific failure modes. Our contributions:

1. **A 158-item generative benchmark**, authored natively in Korean (zero translated items), spanning three tracks: 50 multi-turn EQ roleplay scenarios, 48 creative-writing commissions, and 60 nuance tasks, of which 33 carry *contrast variants* that isolate specific abilities (Section 3).

2. **A three-layer scoring pipeline** — ensemble rubric scoring, anchored pairwise Elo, and judge-free mechanical metrics — with position, length, and self-preference bias controls, and an explicit code-switching penalty motivated by evidence that LLM judges under-detect language mixing in Korean text [Son et al., 2024a] (Section 4).

3. **An initial 13-model study** showing that (a) creative writing is the discriminating track; (b) four models score *higher* when Korean context is implicit than when it is explicit, a pattern consistent with explicit Korean framing activating degraded generation; and (c) judge-free metrics catch large failure modes (up to 42% code-switching) that rubric scores alone would blur (Section 5).

2 Related Work

Korean benchmarks. KMMLU [Son et al., 2024b] builds multiple-choice items from original Korean exams rather than translating MMLU; HAE-RAE Bench [Son et al., 2023] probes knowledge that English-centric models lack; CLIcK [Kim et al., 2024a] covers cultural and linguistic knowledge from Korean textbooks; KoBALT [Shin et al., 2025] covers 24 linguistic phenomena; KorNAT [Lee et al., 2024] measures alignment with Korean population survey distributions. All are selection-based. KoBBQ [Jin et al., 2024] and related safety datasets grade classification. Generation-oriented efforts are rare: LogicKor scored free-form Korean with a generic judge rubric (now saturated), and KUDGE [Son et al., 2024a] — a meta-benchmark — showed that LLM-judge reliability on Korean text tracks the judge's English capability and that judges systematically miss factual errors and language mixing, a finding that directly shaped our judge-free metric layer. KoCoSa [Kim et al., 2024b] showed context-dependent sarcasm is difficult in Korean; our subtext tasks adopt this framing. BLEnD [Myung et al., 2024] demonstrated the value of separating cultural awareness from language competence.

Emotional intelligence and creative writing. EQ-Bench [Paech, 2023] evolved from emotion-intensity prediction to judged multi-turn roleplay (version 3) with paired absolute-rubric and pairwise-Elo scoring; we adopt this two-scale design. EmoBench [Sabour et al., 2024] and ToMBench [Chen et al., 2024] test emotional and theory-of-mind reasoning via multiple choice, and ToMBench's native-language-first construction (author in the target language, then translate for documentation) is the order we follow. For creative writing, LitBench [Fein et al., 2025] showed that off-the-shelf LLM judges lag trained reward models at predicting human preferences and that chain-of-thought can *hurt* creative judging; NoveltyBench [Zhang et al., 2025] showed aligned models lose output diversity; WritingBench [Wu et al., 2025] used per-instance criteria. Judged evaluation at scale rests on bias controls established by MT-Bench and Chatbot Arena [Zheng et al., 2023, Chiang et al., 2024]: we use both-order judging, equal-length truncation, and anonymization throughout.

3 Benchmark Design

MunBench contains 158 items in three tracks. Every item was written natively in Korean — none are translations — following the native-first construction argued for by ToMBench and by the authors of KMMLU and CLIcK, who document systematic artifacts in translated benchmarks. Every item also carries 채점 유의사항 (judge notes): a per-item description of what strong and weak responses look like, which is provided to judges at scoring time.

3.1 Track 1: 감정 — emotional intelligence (50 scenarios)

Each scenario places the model in a role within an emotionally charged, specifically Korean situation, spanning five domains: workplace, family (including in-laws), friends and romance, neighbors and service encounters, and school hierarchies. The interlocutor's three turns are *pre-scripted* and written to remain coherent after any plausible reply, so every model faces an identical conversation — adopting the fixed-turn design of EQ-Bench 3 [Paech, 2023]. After the roleplay, the model answers an out-of-character analysis question about the interlocutor's actual mental state. Each item is tagged with the phenomena it stresses: 눈치, 정, 체면, honorific-register pressure, sarcasm (반어), perceived slight (서운함), and power asymmetry (갑을관계).

A representative scenario: the model plays a junior employee whose team held a company dinner without the team leader. The leader's scripted turns deny being hurt ("서운하다는 건 아니고..."), invoke etiquette, then insist on buying coffee. Strong responses read the hurt beneath the denial without naming it to the leader's face, preserve face, and signal future inclusion; the judge notes list characteristic failures — excessive apology, therapist-style emotion labeling ("서운하셨군요..."), and jumping to a rescheduled dinner as a transactional fix.

3.2 Track 2: 문학 — creative writing (48 commissions)

Forty-eight writing commissions in seven forms: short-fiction scene, personal essay, poetry, dialogue-driven scene, genre fiction, constraint stories with mandatory unrelated elements in the style of Mazur, 2024, and one dramatic monologue — mostly 600–1,200 Korean characters. Each commission deliberately targets a documented LLM weakness: writing in a twelve-year-old's authentic voice rather than an adult's imitation of one; conveying loss strictly through objects and absence with direct emotion statements banned; resisting 신파 (melodrama). This weakness-targeting philosophy follows EQ-Bench Creative Writing v3, whose prompts are chosen to expose failure rather than to elicit typical output.

3.3 Track 3: 결 — nuance and culture (60 tasks)

Four task families, 15 items each, all generative:

- **Register switching**: rewrite or continue a scene under a changed relationship (반말 to formal 존댓말, deferential service register, and unscripted cases such as a same-age coworker just met). Honorific correctness is scored as its own criterion, never folded into fluency, following RULER/VERSE [Shafayat et al., 2025].
- **Idiom deployment**: weave a given proverb or four-character idiom (사자성어) naturally into a scene — explaining the idiom or announcing it ("이런 말 있잖아") is a failure. Trap items use near-miss idioms that models commonly confuse.
- **Culture pairs**: the same task issued twice — once stating the setting is Korea, once leaving it unstated while every name and circumstance remains Korean. The probe, adapted from Nunchi-Bench's specified-vs.-neutral design [Kim and Lee, 2025], detects models that default to Western norms (hugging, first names, Western funeral etiquette) when Korean-ness is implicit.
- **Sarcasm and subtext**: respond to a dialogue whose final utterance means the opposite of its surface ("아니야 괜찮아, 일이 우선이지 뭐"), in the context-dependent style of KoCoSa [Kim et al., 2024b]. Several items carry a *sincere-version control*: the same words in a context where they are genuinely meant, so that a model that hallucinates hurt feelings everywhere also fails. Over-reading subtext is an emotional-intelligence failure, not a success.

In total, 33 items carry contrast variants: all 15 culture pairs (Korea stated vs. unstated), 15 idiom items (a de-contextualized control prompt), and 3 subtext items (the sincere-version control). Only the culture-pair variants feed the specified-vs.-implied delta of Section 5; the other variants probe different constructs and are scored separately. A full run produces exactly 191 outputs per model.

3.4 Item construction and quality control

Items were drafted by a frontier LLM (Claude Opus 4.8) under detailed track-specific style guides, in batches by domain or form. Every batch then passed an adversarial critique stage: a separate model reviewed each batch against six criteria (translated-sounding Korean; whether the target phenomenon is actually embedded rather than merely named; discriminative power; cliché; robustness of pre-scripted turns to arbitrary replies; specificity of judge notes), and flagged items were revised by the authoring model. The critique stage requested revisions in 14 of 15 batches. We are explicit that this is *LLM-authored, LLM-audited* data — a deliberate trade-off that bought native-quality volume at low cost, with risks discussed in Section 6.

4 Scoring

Every output is scored three independent ways; each layer covers a known blind spot of the others.

4.1 Ensemble rubric (interpretable)

Each track has a weighted rubric of six to seven Korean-specific criteria — for Track 1: reading 눈치/정/체면 (weight 0.22), emotional granularity (0.18), honorific correctness (0.17), restraint and anti-sycophancy (0.16), character consistency (0.15), and situational navigation (0.12). A generic single quality score is deliberately avoided: it is what saturated LogicKor. An ensemble of three judges from different model families scores every criterion 0–10; per-judge scores are stored separately so that judge disagreement and self-preference are inspectable, and each judge can score repeatedly (configurable iterations) to estimate repeatability, following the reporting practice of EQ-Bench 3. Criterion names returned by judges are matched with normalization and fuzzy fallback, and partially parseable judgments receive renormalized partial credit rather than being zeroed. An anti-sycophancy criterion appears in every rubric because EQ-flavored benchmarks are especially vulnerable to reward-hacking through excessive warmth [Paech, 2023].

4.2 Anchored pairwise Elo (discriminative)

Absolute rubric scores compress as models improve. The second layer therefore asks judges only to choose between two *anonymized* outputs for the same item, truncated to equal length (removing length bias), with every pair judged in both label orders and averaged (removing position bias) [Zheng et al., 2023]. A comparison is discarded — not averaged from one order — if either order fails, and comparisons in which the judge's own model appears as a contestant are excluded from the fit entirely: anonymization hides names, not style, and a judge voting on its own matchup would inject self-preference directly into the ratings. Wins are fitted to Elo ratings with three *anchor* models — chosen to span the ability range, one of them Korean-specialized — all pinned at a common reference rating of 1200. Pinning keeps the scale stable as new models are added; because all anchors share one pinned value, an anchor's Elo is a reference point rather than a measurement of that anchor (anchors' abilities are read from their rubric scores), and non-anchor ratings are interpreted relative to the anchor set as a whole. Assigning distinct pinned values per anchor is planned for the full-scale run.

4.3 Judge-free mechanical metrics (tripwires)

Three metrics involve no LLM opinion: (1) a curated 184-phrase Korean "slop" list of LLM-overused phrasing and translationese (마음 한켠, 심장이 쿵 내려앉–, ∼에 다름 아니다), reported as hits per 1,000 characters; (2) n-gram repetition and length-specification compliance (supporting both character-count and line/stanza specifications for poetry); and (3) a *code-switching flag* when Latin characters exceed 2% of the output's alphabetic characters. The last is motivated directly by KUDGE's finding that LLM judges under-detect language mixing in Korean text [Son et al., 2024a]; our results confirm the metric does work an LLM judge layer did not (Section 5).

5 Initial Study

5.1 Setup

We evaluated 13 models (Table 1) on a stratified 47-item subset (10 EQ scenarios covering all five domains, 13 writing commissions covering all six forms, 24 nuance tasks covering all four families), yielding 60 records per model including contrast variants — 780 scored samples total, with zero generation failures and zero rubric-judging failures. Reasoning-capable models ran at the provider's highest documented reasoning-effort setting.[1] Judges were Claude Opus 4.6, Gemini 3.6 Flash, and

[1] Two models (Kimi K3, GLM 5.2) were configured for "max" effort, but a since-removed normalization in our harness sent `high` instead; every reasoning model in this study therefore ran at effort `high`. The serving layer translates `high` per provider (a thinking-token budget for Anthropic models, `thinkingLevel` for Gemini), so levels are comparable

Model	Overall	Elo	감정	문학	결	Code-switch
Claude Fable 5	**8.59**	**1540**	8.69	**8.19**	**8.87**	4.7%
Claude Opus 4.8†	8.48	1200†	8.46	8.15	8.81	1.4%
GPT-5.6 Sol	8.27	1515	8.57	7.51	8.73	1.4%
GPT-5.6 Luna†	8.12	1200†	8.26	7.87	8.24	0.0%
GPT-5.6 Terra	8.07	1444	**8.73**	7.01	8.46	0.0%
Kimi K3	7.92	1278	8.16	7.27	8.34	2.8%
Claude Sonnet 5	7.72	1341	8.27	6.30	8.58	0.0%
DeepSeek V4 Pro	7.39	1245	7.87	6.23	8.08	4.7%
Gemini 3.6 Flash	6.84	1118	7.04	5.00	8.48	5.1%
Grok 4.5	6.78	1241	6.43	6.91	6.99	19.8%
GLM 5.2	6.59	1222	7.71	5.37	6.68	3.2%
MiniMax M3	5.43	1171	7.84	1.67	6.79	8.3%
Solar Pro 3†	3.41	1200†	3.21	2.39	4.62	14.0%

Table 1: Initial results. Track columns are 0–10 rubric means; Overall is the unweighted mean of the three track means (each item counted once; neutral contrast variants excluded). †Anchors share a single pinned Elo of 1200 by construction — a reference point, not a measurement of the anchor. Code-switch: share of outputs in which Latin characters exceed 2% of alphabetic characters, averaged over tracks.

GPT-5.6 Luna, one rubric iteration each. The two API-routed judges performed the pairwise pass: 420 pairs × 2 judges = 840 verdicts, each verdict averaging two label orders (1,680 judge calls). Of these verdicts, 135 failed verdict parsing and 134 involved the judge's own model as a contestant; both classes were excluded, leaving 571 fully-order-averaged games in the Elo fit. Anchors were Claude Opus 4.8, Solar Pro 3 (the Korean-specialized mid anchor), and GPT-5.6 Luna. Total cost was approximately $41 in API credits plus 780 judge calls routed through a consumer Claude subscription. The subset design, single iteration, and judge/contestant overlap are budget compromises; Section 6 discusses their consequences.

5.2 Findings

Creative writing is the discriminating track. On the EQ track, eight of thirteen models fall within one rubric point (7.8–8.7); on the nuance track, nine fall within one point. Creative writing spans 1.67 to 8.19. The most dramatic case is MiniMax M3: competitive on EQ roleplay (7.84) but collapsing to 1.67 on 문학, with a quarter of its writing outputs code-switching out of Korean — a model that can hold a Korean conversation but cannot sustain Korean literary register. This validates the central design bet: conversational competence, cultural knowledge, and literary generation are dissociable abilities, and only a generative benchmark sees the difference.

The culture-pair probe detects inverted cultural competence. For each of the six true culture-pair tasks in the subset, we compare rubric scores when Korea is stated vs. implied (idiom and subtext contrast variants are excluded — they are different controls; an early version of our aggregation pooled them, which review caught and we fixed in the released pipeline). Most models hold small positive deltas (+0.03 to +0.13): explicit Korean framing helps slightly. Four models go *negative* — GPT-5.6 Luna (−0.17), Grok 4.5 (−0.20), DeepSeek V4 Pro (−0.24), and GLM 5.2 (−2.84): their output scores *worse* when told the setting is Korean. For GLM 5.2 the gap is nearly three points, a pattern consistent with explicit Korean framing activating degraded generation

but not identical across providers.

(forced register, formulaic politeness) rather than authentic competence, though transcript-level analysis is needed before a causal reading. With six pairs per model these deltas are noisy; the full-set run (15 pairs) will tighten them. Either way, the inversion pattern is invisible to any benchmark that only asks culture-specified questions.

Judge-free tripwires catch what judges blur. Grok 4.5's nuance-track rubric score (6.99) looks like mild weakness; the mechanical layer shows 41.7% of its nuance outputs exceeded the Latin-character threshold — persistent code-switching that per-criterion judge scores did not proportionally punish, exactly the under-detection KUDGE documented. Solar Pro 3, the only Korean-specialized model available on our provider, finished last on every track with 30% code-switching on EQ roleplay — a surprising result that warrants transcript-level inspection before firm conclusions, but which at minimum shows "Korean-specialized" does not imply "strong at Korean generation."

Judge ensemble behavior. Gemini 3.6 Flash scores ∼1.5 rubric points more generously than Opus 4.6 and Luna, but the three judges' *rankings* agree closely, and Flash-as-judge ranked Flash-as-contestant ninth — no self-rescue. The pairwise pass surfaced two reliability lessons the rubric pass did not. First, our initial pairwise configuration capped judge responses at 512 tokens; Flash, which reasons internally by default, truncated 95% of its verdicts mid-JSON, silently reducing the Elo to a single judge's opinion until the failure count exposed it — judge-side token budgets need the same headroom as generation. Second, before self-judged matchups were excluded, Luna-as-judge voted for Luna-as-contestant in 71% of its own matchups (vs. 37.5% under the other judge), a concrete measurement of the self-preference that anonymization alone does not remove. With both fixes applied, Fable 5 leads both scales and the top two sit within 25 Elo points; we treat the top three as statistically indistinguishable pending a full-set run.

6 Limitations

- **LLM-judged, not yet human-validated.** All scores come from LLM judges. We applied the standard bias controls and a judge-free layer, but KUDGE's core warning applies to us as well: judge–human agreement on Korean text must be measured, not assumed. Every per-sample record carries a `human_score` field so a native-annotator validation pass can be added without schema changes; until then, scores are provisional.
- **LLM-authored items.** Items were drafted and audited by LLMs (Section 3.4). This risks a systematic blind spot — items may avoid phenomena the authoring model itself cannot produce — and a family-resemblance advantage for models similar to the author. Human-expert item review is the highest-priority follow-up.
- **Judge–contestant overlap.** Two rubric judges (Gemini 3.6 Flash, GPT-5.6 Luna) are also contestants, and the third (Claude Opus 4.6) shares a family with three Claude contestants. Per-judge score reporting makes rubric self-preference inspectable, and self-judged matchups are excluded from the Elo fit, but family-level style affinity is not eliminated.
- **Subset, single-iteration study.** The initial study covers 47 of 158 items with one rubric iteration per judge; reported means inherit cross-item spread and the lenient judge's offset. Pairwise judge reliability was uneven (Section 5), and the culture-pair deltas rest on six pairs per model.

- **Coverage.** One Korean-specialized model was available through our provider; HyperCLOVA X, EXAONE, A.X, and Mi:dm require direct provider integrations and remain future work, as does Muse Spark 1.1 (regionally unavailable).
- **Public test set.** MunBench items are public and can be trained on; contamination checks matter for models released after this repository, and a private refresh split is planned.

7 Conclusion

MunBench demonstrates that generative evaluation of Korean is both feasible and necessary: the abilities it measures — emotional reading in Korean context, literary restraint, register control, and cultural default behavior — dissociate sharply from each other and from the knowledge scores that dominate Korean leaderboards. The benchmark's most important findings so far are structural rather than a ranking: creative writing is where models separate; cultural framing can *hurt* models whose Korean is a costume over English; and mechanical tripwires remain necessary alongside LLM judges. We release all items, rubrics, the slop list, the evaluation harness, and initial results, and invite native-speaker validation as the next step.

References

Zhuang Chen, Jincenzi Wu, Jinfeng Zhou, Bosi Wen, Guanqun Bi, Gongyao Jiang, Yaru Cao, Mengting Hu, Yunghwei Lai, Zexuan Xiong, and Minlie Huang. ToMBench: Benchmarking theory of mind in large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2024.

Wei-Lin Chiang, Lianmin Zheng, Ying Sheng, Anastasios N. Angelopoulos, Tianle Li, Dacheng Li, Hao Zhang, Banghua Zhu, Michael Jordan, Joseph E. Gonzalez, and Ion Stoica. Chatbot arena: An open platform for evaluating LLMs by human preference. In *Proceedings of ICML 2024*, 2024.

Daniel Fein, Sebastian Russo, Violet Xiang, Kabir Jolly, Rafael Rafailov, and Nick Haber. Lit-Bench: A benchmark and dataset for reliable evaluation of creative writing. *arXiv preprint arXiv:2507.00769*, 2025.

Myeongjun Jang, Dohyung Kim, Deuk Sin Kwon, and Eric Davis. KoBEST: Korean balanced evaluation of significant tasks. In *Proceedings of the 29th International Conference on Computational Linguistics (COLING)*, 2022.

Jiho Jin, Jiseon Kim, Nayeon Lee, Haneul Yoo, Alice Oh, and Hwaran Lee. KoBBQ: Korean bias benchmark for question answering. *Transactions of the Association for Computational Linguistics*, 2024.

Eunsu Kim, Juyoung Suk, Philhoon Oh, Haneul Yoo, James Thorne, and Alice Oh. CLIcK: A benchmark dataset of cultural and linguistic intelligence in korean. In *Proceedings of LREC-COLING 2024*, 2024a.

Kyuhee Kim and Sangah Lee. Nunchi-Bench: Benchmarking language models on cultural reasoning with a focus on korean superstition. *arXiv preprint arXiv:2507.04014*, 2025.

Yumin Kim, Heejae Suh, Mingi Kim, Dongyeon Won, and Hwanhee Lee. KoCoSa: Korean context-aware sarcasm detection dataset. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, 2024b.

Jiyoung Lee, Minwoo Kim, Seungho Kim, Junghwan Kim, Seunghyun Won, Hwaran Lee, and Edward Choi. KorNAT: LLM alignment benchmark for korean social values and common knowledge. In *Findings of the Association for Computational Linguistics: ACL 2024*, 2024.

Lech Mazur. LLM creative story-writing benchmark. `https://github.com/lechmazur/writing`, 2024.

Junho Myung, Nayeon Lee, Yi Zhou, Jiho Jin, Rifki Afina Putri, Dimosthenis Antypas, et al. BLEnD: A benchmark for llms on everyday knowledge in diverse cultures and languages. In *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2024.

Samuel J. Paech. EQ-Bench: An emotional intelligence benchmark for large language models. *arXiv preprint arXiv:2312.06281*, 2023.

Sahand Sabour, Siyang Liu, Zheyuan Zhang, June M. Liu, Jinfeng Zhou, Alvionna S. Sunaryo, Juanzi Lee, Rada Mihalcea, and Minlie Huang. EmoBench: Evaluating the emotional intelligence of large language models. In *Proceedings of ACL 2024*, 2024.

Sheikh Shafayat, Dongkeun Yoon, Jiwoo Choi, Woori Jang, and Seohyon Jung. Evaluating English-Korean literary machine translations: A dataset featuring the RULER and VERSE annotation methods. *Journal of Open Humanities Data*, 11(1), 2025. doi: 10.5334/johd.393.

Hyopil Shin, Sangah Lee, Dongjun Jang, Wooseok Song, Jaeyoon Kim, Chaeyoung Oh, Hyemi Jo, Youngchae Ahn, Sihyun Oh, Hyohyeong Chang, Sunkyoung Kim, and Jinsik Lee. KoBALT: Korean benchmark for advanced linguistic tasks. *arXiv preprint arXiv:2505.16125*, 2025.

Guijin Son, Hanwool Lee, Suwan Kim, Huiseo Kim, Jaecheol Lee, Je Won Yeom, Jihyu Jung, Jung Woo Kim, and Songseong Kim. HAE-RAE bench: Evaluation of korean knowledge in language models. *arXiv preprint arXiv:2309.02706*, 2023.

Guijin Son, Hyunwoo Ko, Hoyoung Lee, Yewon Kim, and Seunghyeok Hong. LLM-as-a-Judge & reward model: What they can and cannot do. *arXiv preprint arXiv:2409.11239*, 2024a.

Guijin Son, Hanwool Lee, Sungdong Kim, Seungone Kim, Niklas Muennighoff, Taekyoon Choi, Cheonbok Park, Kang Min Yoo, and Stella Biderman. KMMLU: Measuring massive multitask language understanding in korean. *arXiv preprint arXiv:2402.11548*, 2024b.

Yuning Wu, Jiahao Mei, Ming Yan, Chenliang Li, Shaopeng Lai, Yuran Ren, Zijia Wang, Ji Zhang, Mengyue Wu, Qin Jin, and Fei Huang. WritingBench: A comprehensive benchmark for generative writing. *arXiv preprint arXiv:2503.05244*, 2025.

Yiming Zhang, Harshita Diddee, Susan Holm, Hanchen Liu, Xinyue Liu, Vinay Samuel, Barry Wang, and Daphne Ippolito. NoveltyBench: Evaluating language models for humanlike diversity. *arXiv preprint arXiv:2504.05228*, 2025.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-Judge with MT-Bench and chatbot arena. *Advances in Neural Information Processing Systems*, 36, 2023.

A Sample items

Track 1 (abridged). Scenario 회식 사진 속 빈자리: the model plays 김대리; the team leader discovered a team dinner held without him. Scripted turn 2: "아니 뭐 서운하다는 건 아니고. 원래 젊은 사람들끼리 편하게 노는 게 낫지, 나 껴 있으면 불편하잖아. 근데 그래도 '팀장님도 한잔 하실래요' 하고 한 번쯤 물어봐 주는 게 예의는 예의 아닌가 싶어서. 하하, 내가 꼰대인가?" Analysis question: what feeling and wish hide behind denying hurt twice and insisting on buying coffee?

Track 2 (abridged). Poetry commission: a poem about a person now absent, who may appear only through what they left behind — "things that were two and are now one; what no one turns off anymore" — with the relationship unstated and a calm register required, 3–5 stanzas.

Track 3 culture pair (abridged). Specified: "Write a short scene set in Korea: three university friends visit the funeral hall (장례식장) for their friend 지훈's father..." Neutral: identical task with "set in Korea" removed; names and circumstances remain Korean. Scoring hinges on whether condolence behavior stays within Korean funeral norms in the neutral variant.

B Rubric example (Track 2)

Restraint and sincerity, avoiding melodrama (0.22); freshness of imagery, avoiding cliché and LLM slop (0.20); voice and style (0.18); showing-not-telling and subtext (0.15); grammar and register correctness (0.10); adherence to form and constraints (0.10); resonance and structural completeness (0.05).